

Reply to Reviewer#2

I. In the main idea, it is not clear the sampling proportion between the data from known and unknown classes. How does this proportion affect the subsequent training?

Response:

As you mentioned, it is essential to indicate the affection of proportion between the data from known and unknown classes in the setting of open set domain adaptation. However, in real scenarios, we can not manually set the sampling proportion due to the lack of target label information. During the training phase, ThDAN is confined to randomly sample training data from the target domain. So we conduct experiments to implicitly change the sampling proportion by varying the number of unknown data and the number of unknown classes. All experiments are conducted on Office-31 datasets. The following figures show the classification results.

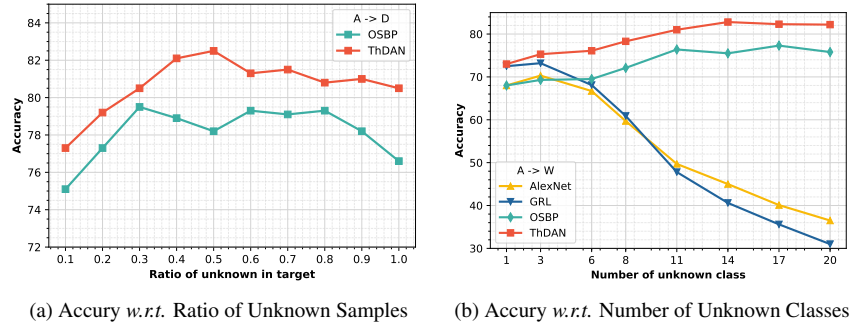


Fig. 1: (a): The prediction accuracy of our method when we changed the ratio of unknown samples in the adaptation task $A \rightarrow D$. (b): The prediction accuracy of our method when we changed the number of unknown classes in the adaptation task $A \rightarrow W$.

We first fix the number of unknown classes to 10, and then change the sample number of each unknown class. **Ratio 0.1** of sample number indicates that only one-tenth samples of each unknown class are used for training and testing, while the others are discarded from the dataset. Fig.1a shows the result. When the ratio is low, the model can hardly identify samples from the unknown class due to the lack of training samples. On the contrary, when the ratio is high, the negative effect of the unknown class will deteriorate the accuracy of categorizing known classes. Therefore the performance of the model is defective around the endpoint of the interval.

We further vary the number of unknown classes in Fig.1b. As the number of unknown classes increases, the performance of the methods proposed for close set domain adaptation will decrease significantly. That is because they cannot alleviate the negative transfer brought by the unknown class. The model proposed for open set domain adaptation would gain performance improvement by correctly rejecting samples of the unknown class. This is because it is easier to reject samples as the “unknown” than categorize samples from the diverse known classes. The accuracy tends to decay since too many unknown classes will deteriorate the process of knowledge transferring for the known classes.

The experiments show that the sampling proportion between the data from known and unknown classes could substantially affect model training for both unknown and known classes. The low proportion of samples from unknown classes will affect the ability to identify samples from the unknown class. On the other hand, the high proportion of that will affect the ability to classify samples from known classes.

II. From their reported results in Fig.7, the adjust threshold of γ_0 is not sensitive to the accuracy of training, it is a bit strange. Authors are encouraged to carefully check their claims by using more detailed reasons or more diverse data sets.

Response:

Thanks for your suggestions. In our model, the threshold is calculated by averaging the transferability score of a batch of source samples and the predefined γ as follows,


$$\beta(X_s, \gamma) = \mathbb{E}_{x \in X_s} w(x) - \gamma. \quad (\text{R2.1})$$

Then we dynamically increase γ from 0 to γ_0 so that more samples can be selected for domain adversarial training. The changing function of γ is as follows,

$$\gamma = \begin{cases} 0 & , n \in N_1 \\ \gamma_0 \times \sigma(n) & , n \in N_2 \end{cases}. \quad (\text{R2.2})$$

Here σ is a *monotonically increasing function* with an upper bound of 1. N_1 and N_2 denote different training stage.

In experiments, we verify that the model is insensitivity to the value of γ_0 on **6 tasks of Office-31** and **12 tasks of Office-Home**. We include Fig.7 of the original paper as follows, which shows the average prediction accuracy on these tasks when γ_0 changes,



(a) Accury w.r.t. value of γ_0

Fig. 2: The prediction accuracy of our method when we change the value of σ_0 on dataset Office-31 and Office-Home.

Compared with the baseline model OSBP, our model can deliver plausible results on both Office-31 and Office-Home datasets when γ_0 in the range of $[0.06, 0.5]$. That is because when we gradually increase the γ during training, we also gradually decrease the learning rate. As a result, the transferable samples selected in the early stage would “dominate” the model training, because they engage in training sooner, and the gradient produced by them can update the model with a larger learning rate. As the γ approaches γ_0 , the learning rate is relatively low, and the previously selected samples would still be selected for training; therefore, the newcomers have little impact on model performance. This directly reflects on the insensitivity of value of γ_0 to the accuracy of training.

The above explanation has been added in the revised manuscript. Please refer to Section 5.3.6 in the revised manuscript for details.

1. The abstract needs to be redesigned. Open set domain adaptation (OSDA) should be a scenario. Authors claim to solve the OSDA problem. What is the problem?? This is unclear. It seems the author should firstly explain the OSDA setting and then present which issue they want to address. The first sentence of the abstract is a bit far away from their focus. More seriously, the introduction don not connect the abstract. For example, the beginning of the abstract explains the unsupervised domain adaptation, but the beginning of the introduction suddenly use deep learning to start this paper without unsupervised domain adaptation. Authors are encouraged to carefully reorganize their work, where the introduction must connect the abstract tightly. One sentence in the abstract should connect one logic in the introduction section.

Response:

Thanks for your comments, the suggestions are of great help to improve our work. In the revised manuscript, we redesigned the **abstract** and **introduction** for tighter logical connections.

The **abstract** has being redesigned as follows,

In recent years, many unsupervised domain adaptation (UDA) methods have been proposed to tackle the domain shift problem. Most existing UDA methods are derived for Close Set Domain Adaptation (CSDA) in which source and target domains are assumed to share the same label space. However, target domain may contain unknown class different from the known ones in the source domain in practice, i.e., Open Set Domain Adaptation (OSDA). Due to the presence of unknown class, aligning the whole distribution of the source and target domain for OSDA as in the previous methods will lead to negative transfer. Existing methods developed for OSDA attempt to assign smaller weights to target samples of unknown class. Despite promising performance achieved by existing methods, the samples of the unknown class are still used for distribution alignment, which make the model suffer from the risk of negative transfer. Instead of reweighting, this paper presents a novel method namely Thresholded Domain Adversarial Network (ThDAN), which progressively selects transferable target samples for distribution alignment. Based on the fact that samples from the known classes must be more transferable than target samples of the unknown one, we derive a criterion to quantify the transferability by constructing classifiers to categorize known classes and to discriminate unknown class. In ThDAN, an adaptive threshold is calculated by averaging transferability scores of source domain samples to select target samples for training. The threshold is tweaked progressively during the training process so that more and more target samples from the known classes can be correctly selected for adversarial training. Extensive experiments show that the proposed method outperforms state-of-the-art domain adaptation and open set recognition approaches on benchmarks.

According to the redesigned abstract, we have modified the **introduction** to make it connect to the abstract tightly. The following Table 1 shows the logical connection between each paragraph of the introduction and each sentence of the abstract.

Table 1: The logical connection between paragraphs of the introduction and sentences of the abstract.

Paragraph of Introduction	Sentence of Abstract
# 1	<i>In recent years, many unsupervised domain adaptation (UDA) methods have been proposed to tackle the domain shift problem.</i>
# 2	<i>Most existing UDA methods are derived for Close Set Domain Adaptation (CSDA) in which source and target domains are assumed to share the same label space. However, target domain may contain unknown class different from the known ones in the source domain in practice, i.e., Open Set Domain Adaptation (OSDA).</i>
# 3	<i>Existing methods developed for OSDA attempt to assign smaller weights to target samples of unknown class. Despite promising performance achieved by existing methods, the samples of the unknown class are still used for training, which make the model suffer from the risk of negative transfer.</i>
# 4	<i>Instead of reweighting, this paper presents a novel method namely Thresholded Domain Adversarial Network (ThDAN), which progressively selects transferable target samples for distribution alignment.</i>

2. What is “transferable”? How to evaluate it? Any definition to support this term?

Response:

Unfortunately, there is no formal definition or mathematical formula to define “transferable” in the context of domain adaptation. Generally, the word “transferable” is used to describe features. Here we quote from [1]:

the transferable features are the features that generalize well to novel tasks for domain adaption

While the later work [2] also uses “transferable” to describe samples that contribute to the transfer task of domain adaptation.

In the scenario of open set domain adaptation, we consider the target samples from the known classes are transferable samples since they can boost the classification performance of the model. On the contrary, the target samples from the unknown class are considered as untransferable samples, that is because aligning distribution with them will incur negative transfer.

In this work, we evaluate transferability based on two observations:

Target samples from the known classes are able to confuse G_d . Here G_d is the domain discriminator that gives the probability of being target samples. For a target sample, if $G_d(z)$ approaches to 1, then the sample has a high probability of coming from the unknown class. That is because the unknown class is only included in the target domain and can be almost perfectly discriminated from the source samples. On the other hand, if $G_d(z)$ approaches to 0, then the sample is more likely from the known classes that shared by domains. Therefore the transferable target samples are able to confuse G_d to label them as the source samples. Then we can define the transferability $w_d(x)$ as inversely related to $G_d(z)$,

$$w_d(x) = 1 - G_d(z). \quad (\text{R2.3})$$

Target samples from the known classes can be categorized by $G_{c,known}$. Here $G_{c,known}$ is the classifier that gives the probability distribution among known classes. Due to the overlapping in the marginal distributions across domains, target samples from the known classes could be categorized correctly by the classifier $G_{c,known}$ that trained on the source samples. This leads to low entropy (high transferability) for the samples from the known classes. For the samples from the unknown class, because they cannot be classified into one of the K known classes, the uncertainty of the prediction measured by the entropy is large. Therefore the transferability $w_c(x)$ can be defined as inversely related to the *normalized* classification entropy H ,

$$w_c(x) = 1 - H(G_{c,known}(z)). \quad (\text{R2.4})$$

Since the domain discriminate G_d in Eq.(R2.3) and the classifier $G_{c,known}$ in Eq.(R2.4) work independently, we can unify the transferability criterion as,

$$w(x) = 1 - G_d(z) \cdot H(G_{c,known}(z)). \quad (\text{R2.5})$$

The experiments in Section 4 show that the transferability criterion of Eq.(R2.5) works well for our model to select target samples from the known classes.

3. In Fig.1, it is unclear why (d) must outperform the other methods.

Response:

Thanks for your comments. Fig.1 in the original manuscript is as follows,

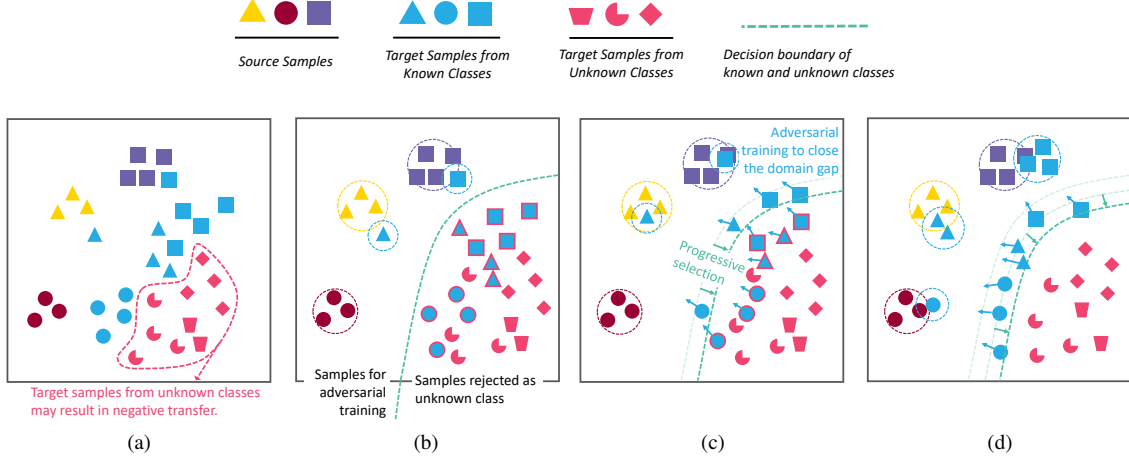


Fig. 3

As a matter of fact, Fig.3 illustrates the general idea of the proposed method. To make a clearer presentation for Fig.3, we modify the caption in the revised manuscript as follows,

The general idea of the proposed Thresholded Domain Adversarial Network (ThDAN). (a): Training samples for open set domain adaptation. (b): ThDAN builds a decision boundary based on the transferability threshold to separate known and unknown classes. (c): ThDAN aligns the distributions of the source and the samples selected as known classes. Then ThDAN tweaks the transferability threshold to collect more samples of known classes to enhance distribution alignment. (d): ThDAN progressively selects more samples of the known classes for training. All the unselected samples are rejected as “unknown”.

4. What is “domain-invariant features” ?

Response:

In domain-invariant feature space, the source and target domains have the same (or similar) marginal distributions, and the posterior distributions of the labels are the same across domains too [3]. Hence, a classifier trained on the labeled source domain would likely perform well on the target domain.

Specifically, the domain-invariant features can be obtained by performing *domain adversarial training* [4, 5, 6]. The domain adversarial training is a minmax game: a domain discriminator is trained to separate the feature representation of the source domain from the one of the target domain, at the same time, feature generators are trained to deceive the domain discriminator. In this work, we adopt the training scheme proposed in [6] to enable the domain discriminator to identify samples of unknown class for open set domain adaptation. Formally, the training procedure can be written as,

$$\min_{G_f^s} \max_{G_d} \mathcal{L}(G_d, G_f^s, G_f^t) = \mathbb{E}_{x \sim p_t(x)} \left[\log \left(G_d \left(G_f^t(x) \right) \right) \right] + \mathbb{E}_{x \sim p_s(x)} \left[\log \left(1 - G_d \left(G_f^s(x) \right) \right) \right]. \quad (\text{R2.6})$$

G_f^s and G_f^t are the feature extractors for source and target samples, which share weights as in [6]. G_d aims to label source samples as 0 and target samples as 1. The optimization of Eq.(R2.6) makes G_f to generate features that can not be discriminated by G_d (i.e., domain-invariant features). Therefore the classifier G_c that trained on the source domain can perform well for target samples by leveraging features extracted by G_f .

5. In the related work, CSDA does not have references.

Response:

Thanks for your suggestion. CSDA is referred to Close Set Domain Adaptation and we have add references [7, 8, 9] for the methods for CSDA in the revised manuscript.

6. In section 2.1, "Another way to solve domain adaptation" is not suitable. Domain adaptation is a setting in detailed ML tasks or a probability distribution issue. Authors can obtain more information from S Ben-David's paper.

Response:

Thanks for your comment. We agree that the phrase "Another way to solve domain adaptation" is not precise. This phrase is revised as "Another approach for domain adaptation". Similar unprecise wordings have been modified in the revised manuscript. Also, we have carefully read Ben-David's paper [7] to get a better understanding of domain adaptation and has cited this paper in the revised manuscript.

7. Section 2.2 missed some new work from open-set domain adaptation.

Response:

Thanks for your suggestions. We add some new work [10, 11] of open set domain adaptation in Section 2.2, they are respectively from ICML2019 and CVPR2019.

Furthermore, we implement *Factorized Representations For Open Set Domain Adaptation (FRFOSDA)* model proposed in [10] in experiments on **Office-Hone** and **Office31** for comparisons. Please refer to Section 4 for more details.

8. In Section 3, authors missed the definition on C_t . Is it the label space of target domain?

Response:

Yes, C_t is the label space of target domain. In the revised manuscript, the definition on C_t has been added in section 3.1 *open set domain adaptation* as follows,

C_s and C_t respectively denote the label space of the source and target domain. In the setting of open set domain adaptation, the label space of target domain contains the label space of source domain, *i.e.*, $C_s \subset C_t$. We refer to classes from C_s as the known classes and classes from $C_t \setminus C_s$ as the unknown class.

9. In Section 3.1, what is “untransferable ones”?

Response:

This question is related to Question 2. In the setting of open set domain adaptation, the untransferable ones refer to the samples of unknown class in the target domain. The proposed method selects untransferable samples by evaluating transferability as follows,

$$X_t^u = \{x | w(x) < \beta, x \in X_t\}. \quad (\text{R2.7})$$

The w is transferability calculator in Eq.(R2.5), and β is transferability threshold calculated by averaging and tweaking transferability scores of source samples. The selection procedure indicated in Eq.(R2.7) shows that target samples X_t whose transferability smaller than the threshold will be selected as untransferable target samples X_t^u .

10. In Eq.(2), more strong reasons need to be presented to explain the G_f and G_d . Why the use such settings to define them? Only based other’s work?

Response:

Thanks for your comments. The reason why we use the setting “ G_f and G_d can be considered as the generative network and discriminate network of GAN respectively” to define G_f and G_d in Eq.(2) is that the *domain adversarial training* between G_f and G_d is similar to the original training procedure of *Generative Adversarial Nets (GAN)* [12]. In GAN, the discriminate network D and the generative network G are trained with the following objective,

$$\begin{aligned} \min_G \max_D V(D, G) = & \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] \\ & + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]. \end{aligned} \quad (\text{R2.8})$$

In the training of GAN of Eq.(R2.8), D is trained to separate the feature representation of the true images and fake images, and G is trained to generate fake images that can confuse D . While in the domain adversarial training of Eq.(R2.6), G_d is trained to separate the feature representation of the source domain from the target domain, and G_f is trained to generated target features that are able to confuse G_d . The training procedure for D and G is very similar to the domain adversarial training for G_d and G_f . Therefore we can define G_f and G_d by the setting of GAN.

We think this setting is not convincing enough to explain why Eq.(2) is the optimal G_d , and have removed it in the revised manuscript. We have added more formal proof for Eq.(R2.9) in the revised manuscript, please refer to the reply of the next question for details.

11. In the following, why does G_d converge to the optimal?

Response:

Eq.(2) is as follows,

$$G_d^*(z) = \frac{p_t(z)}{p_s(z) + p_t(z)}. \quad (\text{R2.9})$$

Here z is the deep feature representation of a sample, *i.e.*, $z = G_f(x)$. We give the proof that Eq.(R2.9) is the optimal G_d for Eq.(R2.6) as follows,

Proof. For any G_f^s and G_f^t , we train G_d to maximize Eq.(R2.6):

$$\begin{aligned}\max_{G_d} \mathcal{L}(G_d, G_f^s, G_f^t) &= \int_x p_t(x) \log(G_d(G_f^t(x))) + p_s(x) \log(1 - G_d(G_f^s(x))) dx. \\ &= \int_z p_t(z) \log(G_d(z)) + p_s(z) \log(1 - G_d(z)) dz.\end{aligned}\tag{R2.10}$$

We take the partial differential of the objective Eq.(R2.10) with respect to G_d , and apply the Leibniz rule to exchange the order of differentiation and integration to achieve optimal G_d in $[0, 1]$ at Eq.(R2.9). \square

The proof has been added in Section 3.2 of the revised manuscript.

12. In Eq.(2), how to define p_s and p_t ?

Response:

Eq.(2) is as follows,

$$G_d^*(z) = \frac{p_t(z)}{p_s(z) + p_t(z)}.\tag{R2.11}$$

Here z is the deep feature representation of a sample, *i.e.*, $z = G_f(x)$. As stated in *Section 3.1 open set domain adaptation* of the revised manuscript, $p_s(z)$ and $p_t(z)$ are probability distributions of feature representations in source and target domains respectively, as defined in [13].

13. In Eq.(9), what is X_t^k ?

Response:

Thanks for your comments. The Eq.(9) in the original manuscript is as follows,

$$\begin{aligned}X_t^k &= \{x | w(x) \geq \beta, x \in X_t\}, \\ X_t^u &= \{x | w(x) < \beta, x \in X_t\}.\end{aligned}\tag{R2.12}$$

Here X_t^k denotes the transferable target samples that are selected as the samples from the known classes. And X_t^u denotes the untransferable target samples that are regarded as the samples from the unknown class. Specifically, in Eq.(R2.12), β is a transferability threshold, w is the transferability calculator. For a batch of target training data X_t , we use the threshold β to split it into two parts, X_t^k and X_t^u . For target samples in X_t^k , we train the model to align distribution with the source samples. For target samples in X_t^u , we train the G_d to identify them as “unknown” samples.

We added Section 4.1 in the revised manuscript to clarify the proposed methods and equations, in which X_t^k and X_t^u are explained.

14. A same controversial issue appears again: Sample selection is the key step for this model to address open set domain adaptation.

Response:

Thanks for pointing out the mistakes. This sentence has been modified as “The sample selection algorithm will substantially affect the performance of ThDAN for open set domain adaptation”

Similar mistakes have been corrected in the revised manuscript.

15. In the experiments, why do you use different setting to begin the experiments in Section 4.2.1? Authors are encouraged to explain the inherent reasons.

Response:

Thanks for your suggestions. We want to use different settings to perform ablation studies to evaluate the efficacy of the threshold tweaking techniques proposed in Section 4.5. The ablation studies on Office-31, OfficeHome and VisDA datasets involve 3 kinds of ThDAN variants. (1) *ThDAN-m-dy* is the original setting which calculates transferability threshold based on mini-batch samples and a fixed γ . (2) *ThDAN-dy* is the variant applies exponential moving average to update transferability threshold. (3) *ThDAN* further takes advantage of dynamic γ to tweak transferability threshold.

In order to better understand how the proposed method performs under ablation settings, we add Section 5.3.1 *Ablation Study* in the revised manuscript. Please refer to the revised manuscript for details.

16. A same issue appear again in the settings of Section 4.2.2.

Response:

Thanks for your comments. This question is related to the previous question. In the revised manuscript, the ablation study in the original Section 4.2.2 has been reorganized in Section 5.3.1 *Ablation Study*.

17. The results analysis in Section 4.2.3 is unclear. More detailed reasons need to be present to support your conclusion.

Response:

The results of Section 4.2.3 are as follows,

Table 2: Accuracy (%) of each method with 6 shared class on VisDA.

Method	class-wise accuracy on VisDA								
	Bicycle	Bus	Car	Mcycle	Train	Truck	Unkown	Avg.All	Avg.Known
OSVM [ECCV, 2014]	4.8	45.0	44.2	43.5	59.0	10.5	57.4	37.8	34.5
MMD [NIPS, 2007]	0.2	30.9	49.1	54.8	56.1	8.1	61.3	37.2	33.2
GRL [JMLR, 2016]	9.1	50.5	53.9	79.8	69.0	8.1	42.5	44.7	45.1
OSBP [ECCV, 2018]	48.0	67.4	39.2	80.2	69.4	24.9	80.3	58.5	54.8
ThDAN-dy	52.6	69.5	58.8	83.1	72.3	12.3	84.5	61.8	58.1
ThDAN	55.2	70.8	61.4	85.6	74.8	10.0	88.8	63.8	59.6

In the experiments of VisDA, the training split is used as the source domain and validation one as the target domain. We choose 6 categories, i.e., bicycle, bus, car, motorcycle, train and truck as the known classes, and the other 6 categories as the unknown class. As Table. 2 shows the classification results, where *Avg.All* and *Avg.Known* indicate the accuracy averaged over all classes and known classes. For the classification of known classes, ThDAN can exceed other methods in almost every class and in the average, which means our method can effetely transfer knowledge for the classification task. This is because ThDAN can avoid the negative transfer by only selecting transferable target samples for domain adversarial training. Also, our model improves the classification accuracy for the unknown class

by a big margin. This further verifies the efficiency of the sample selection algorithm, since ThDAN will reject the unselected samples as the unknown class.

The analysis mentioned above has been added to the corresponding section in the revised manuscript.

18. Please carefully check the claim in Fig.7.

Response:

Thanks for your reminding. The Fig.7 in the original manuscript is as follows,




Fig. 4

In the revised manuscript, the claim of Fig.4 is modified for clearer representation. The rewritten claim has been stated in the previous questions:

- In **Question I**, we modify the claim of Fig.4a and 4b by adding more analysis to explain how the performance changes with different sampling proportions between the data from known and unknown classes.
- In **Question II**, we modify the claim of Fig.4c by providing a more detailed analysis on why the proposed method is insensitive to the setting of γ_0 .

19. Some equations missed “,” or “.” at their ends.

Response:

Thanks for your reminding. The equations in the revised manuscript have been corrected.

References

- [1] M. Long, Y. Cao, J. Wang, M. I. Jordan, Learning transferable features with deep adaptation networks, arXiv preprint arXiv:1502.02791 (2015).
- [2] X. Wang, L. Li, W. Ye, M. Long, J. Wang, Transferable attention for domain adaptation, in: Thirty-Third AAAI Conference on Artificial Intelligence, 2019.
- [3] E. Tzeng, J. Hoffman, N. Zhang, K. Saenko, T. Darrell, Deep domain confusion: Maximizing for domain invariance, arXiv preprint arXiv:1412.3474 (2014).
- [4] Y. Ganin, E. Ustinova, H. Ajakan, P. Germain, H. Larochelle, F. Laviolette, M. Marchand, V. Lempitsky, Domain-adversarial training of neural networks, The Journal of Machine Learning Research 17 (1) (2016) 2096–2030.

- [5] E. Tzeng, J. Hoffman, K. Saenko, T. Darrell, Adversarial discriminative domain adaptation, in: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2017, pp. 7167–7176.
- [6] K. Saito, S. Yamamoto, Y. Ushiku, T. Harada, Open set domain adaptation by backpropagation, in: Proceedings of the European Conference on Computer Vision, 2018, pp. 153–168.
- [7] S. Ben-David, J. Blitzer, K. Crammer, A. Kulesza, F. Pereira, J. W. Vaughan, A theory of learning from different domains, *Machine learning* 79 (1-2) (2010) 151–175.
- [8] M. Wang, W. Deng, Deep visual domain adaptation: A survey, *Neurocomputing* 312 (2018) 135–153.
- [9] S. J. Pan, Q. Yang, A survey on transfer learning, *IEEE Transactions on Knowledge and Data Engineering* 22 (10) (2010) 1345–1359.
- [10] M. Baktashmotlagh, M. Faraki, T. Drummond, M. Salzmann, Learning factorized representations for open-set domain adaptation, in: International Conference on Learning Representations, 2019.
- [11] H. Liu, Z. Cao, M. Long, J. Wang, Q. Yang, Separate to adapt: Open set domain adaptation via progressive separation, in: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2019, pp. 2927–2936.
- [12] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, Y. Bengio, Generative adversarial nets, in: Advances in Neural Information Processing Systems, 2014, pp. 2672–2680.
- [13] J. Zhang, Z. Ding, W. Li, P. Ogunbona, Importance weighted adversarial nets for partial domain adaptation, in: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2018, pp. 8156–8164.